Padmanaba Srinivasan

padmanabasrinivasan@gmail.com padmanaba-srinivasan
padmanabasrinivasan.github.io

Experience

Mar. 2025 – Present **Research Scientist @ Meta**
Applied ML Research in Ads Revenue and Ranking

Jan. 2021 – Oct. 2021 **Researcher (intern) @ Infosys**
Computer vision, imitation learning, and reinforcement learning applied to tennis with the Infosys Tennis Platform team

Apr. 2019 – Sept. 2019 **Software Engineer (intern) @ Credit Suisse**
NLP & Chatbot assistants for development and trading

Jun. 2018 – Sept. 2018 **Software Engineer (intern) @ GCHQ**
Cybersecurity & Communications

Education

2020 – 2025 **PhD, Imperial College London** Computing (Machine Learning)
Thesis title: *Offline Reinforcement Learning: In Pursuit of Perfect Policies from Imperfect Data*
Advisor: William J. Knottenbelt

2016 – 2020 **MEng, Imperial College London** Electronic and Information Engineering
First-class honours
Thesis title: *Machine Learning for the Analysis and Prediction of Film Performance*

Awards

2023 **Finalist, Sloan Sports Analytics Conference Research Paper Competition**
Research Paper selected as finalist, oral presentation.

2020 **Winner, Citadel European Datathon**
Developed methodology to identify gentrifying areas in NYC.

Distinguished Project Award
Awarded by Imperial's Department of Computing for MEng thesis.

Publications (selected)

1. P. Srinivasan and W. Knottenbelt, "Behaviour Preference Regression for Offline Reinforcement Learning," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 39, 2025, pp. 20 575–20 583.
2. P. Srinivasan and W. Knottenbelt, "Offline Model-Based Reinforcement Learning with Anti-Exploration," in *Proceedings of the 27th European Conference on Artificial Intelligence*, 2024.
3. P. Srinivasan and W. Knottenbelt, "Offline Reinforcement Learning with Behavioral Supervisor Tuning," in *Proceedings of the Thirty-Third International Joint Conference on Artificial Intelligence, IJCAI-24*, International Joint Conferences on Artificial Intelligence Organization, Aug. 2024, pp. 4929–4937. DOI: 10.24963/ijcai.2024/545.
4. P. Srinivasan and W. J. Knottenbelt, "SpOiLer: Offline Reinforcement Learning using Scaled Penalties," in *6th Annual Learning for Dynamics & Control Conference*, PMLR, 2024, pp. 825–838, ISBN: 2640-3498.